

第十一章 空间网络模型与查询

陶煜波

计算机科学与技术学院

几何对象模型

- 空间数据模型分类

- 矢量模型

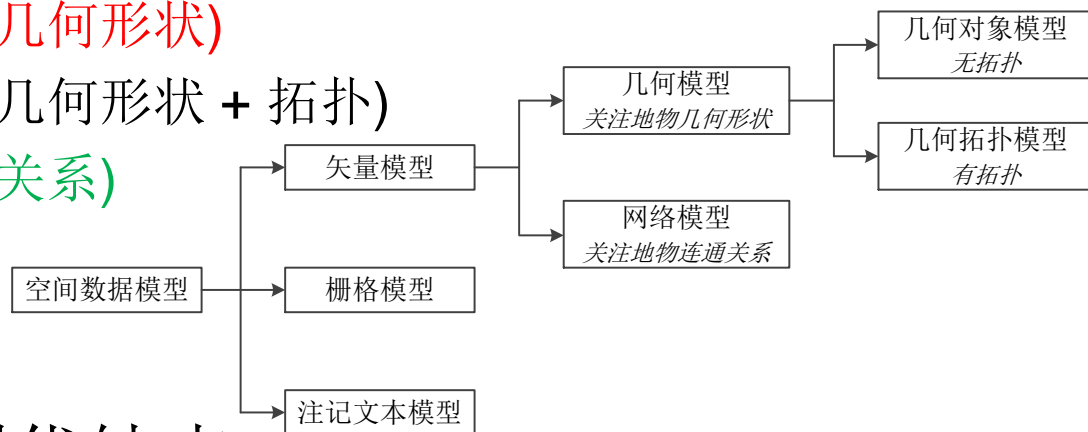
- 几何对象模型 (地物几何形状)

- 几何拓扑模型 (地物几何形状 + 拓扑)

- 网络模型 (地物连通关系)

- 栅格模型

- 注记文本模型



- 矢量模型和栅格模型优缺点

- 概念模型 → 逻辑模型 → 物理模型

几何对象模型

- 对象关系数据库：数据+方法（C++中的类）

- 几何对象层次关系

- 坐标维数和几何维数
- 边界、内部、外部
- 九交矩阵

- 几何对象方法

- 常规方法 (12种)

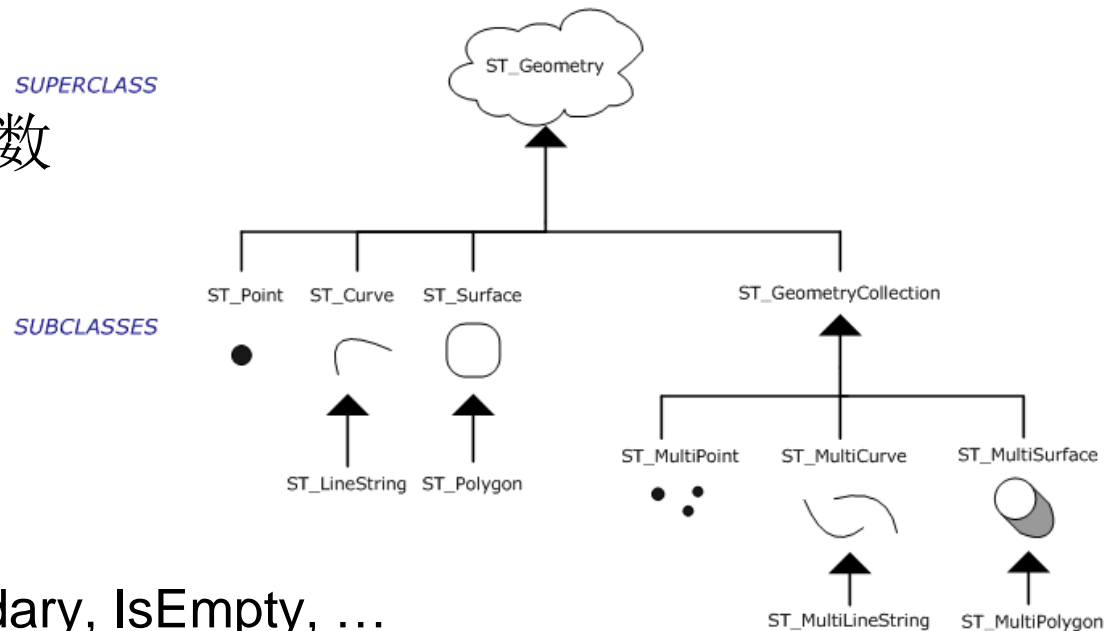
- Dimension, Boundary, IsEmpty, ...

- 常规GIS分析方法 (7种)

- Distance, Buffer, ConvexHull, Intersection, Union, Difference, SymDifference

- 空间查询方法 (11种)

- 包含于(within): 若 $a \cap b = a$, 且 $I(a) \cap E(b) = \emptyset$, 则a包含于b内



几何对象模型

- 逻辑模型

- 基于预定义数据类型的实现

- numeric和BLOB

- 基于扩展几何类型的实现

- Geometry类

- 表模式

- 系统表

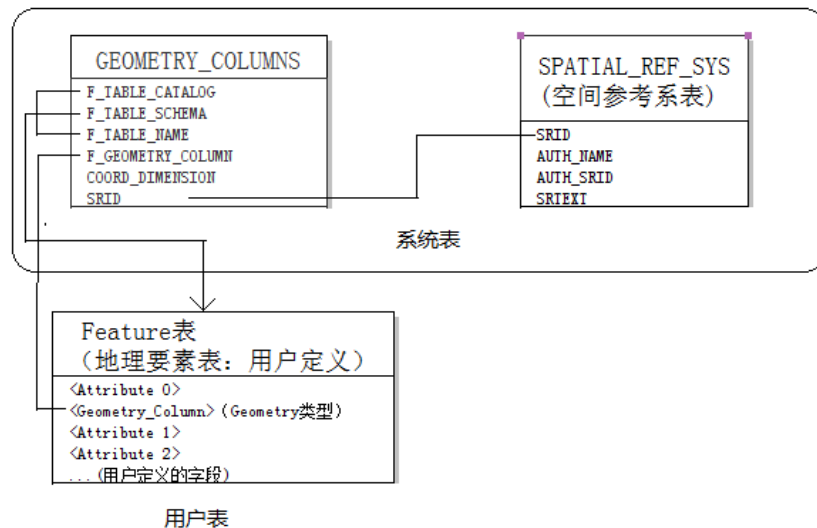
- GEOMETRY_COLUMNS和SPATIAL_REF_SYS

- 用户表

- Feature和Geometry

- 物理模型

- WKB和WKT



基于扩展Geometry类型的要素表模式

空间数据库

- 空间数据库 = 对象关系/关系数据库 + 空间扩展
 - Oracle + Oracle Spatial
 - SQL Server + SQL Server Spatial
 - PostgreSQL + **PostGIS**
 - MySQL + MySQL Spatial
 - SQLite + SQLite Spatialite
- PostGIS
 - 提供了空间数据类型、空间函数和空间索引
 - Geometry (Point/Line/Polygon/Multixxx, 空间参考系)
 - ST_XXX
 - GiST

应用举例：打车软件

- 通常，数据库保留所有时间上的信息，即
 - Taxi(ID, driverID,, status, pos(Point, 4326), time)
 - User(ID, name,, pos(Point, 4326), time)
 - Road(ID, name,, line(LineString, 4326))
- 应如何修改上述SQL语句？
- 出租车(ID = B)附近1公里内的乘客叫车？
 - Select T.ID, T.position
From Taxi T, User U
Where T.ID = B and ST_Distance(T.pos, U.pos) < 1000
and T.time
and U.time
— 这样的实现方式合理吗？

第十一章 空间网络模型与查询

- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Navigation Systems

- Historical
 - Navigation is a core human activity for ages!
 - Trade-routes, Routes for Armed-Forces
- Recent Consumer Platforms
 - Devices: Phone Apps, In-vehicle, “GPS”, ...
 - WWW: Google Maps, MapQuest, ...
- Services
 - Display map around current location
 - Compute the shortest route to a destination
 - Help drivers follow selected route



Location Based Services

- Location: Where am I?
 - Geo-code: Place Name (or Street Address) → <latitude, longitude>
 - Reverse Geo-code: <latitude, longitude> → Place Name
- Directory: What is around me?
 - Where is the nearest Clinic? Restaurant? Taxi?
 - List all Banks within 1 mile
- Routes: How do I get there?
 - What is the shortest path to get there?
 - ...

基于位置的服务

- 休闲娱乐型
 - 签到模式
 - 大富翁游戏模式
 - 生活服务型
 - 周边生活服务的搜索
 - 与旅游的结合
 - 会员卡与票务模式
 - 社交型
 - 地点交友，即时通讯
 - 以地理位置为基础的小型社区
 - 商业型
 - LBS+团购；优惠信息推送服务；店内模式
- <http://blog.csdn.net/superjunjin/article/details/7818378>

Spatial Networks & Modern Society

- Transportation, Energy, Water, Communications,

...



Quiz 1

- Which of the following is not an application of spatial networks?
 - a) Navigation system
 - b) Geocoding
 - c) Reverse geocoding
 - d) Convex hull of a country

Limitations of Spatial Querying

- OGIS Simple Feature Types
 - Supports Geometry (e.g., Points, LineStrings, Polygons, ...)
 - However, lack **Graphs** data type, **shortest_path** operator
- Traditional SQL
 - Supports select, project, join, statistics
 - Lacked transitive closure, e.g., **network analysis** (next slide)
 - SQL3 added **recursion** & **transitive closure**

Spatial Network Analysis

- Route (A start-point, Destination(s))
 - What is the shortest path to get there?
 - What is the shortest path to cover a set of destinations?
- Allocation (A set of service centers, A set of customers)
 - Assign customers to nearest service centers
 - Map service area for each service center
- Site Selection (A set of customers, Number of new service centers)
 - What are best locations for new service centers?

Quiz 2


- Which of the following is not included in OGIS simple feature types?
 - a) Graph
 - b) Point
 - c) Line String
 - d) Polygon

第十一章 空间网络模型与查询

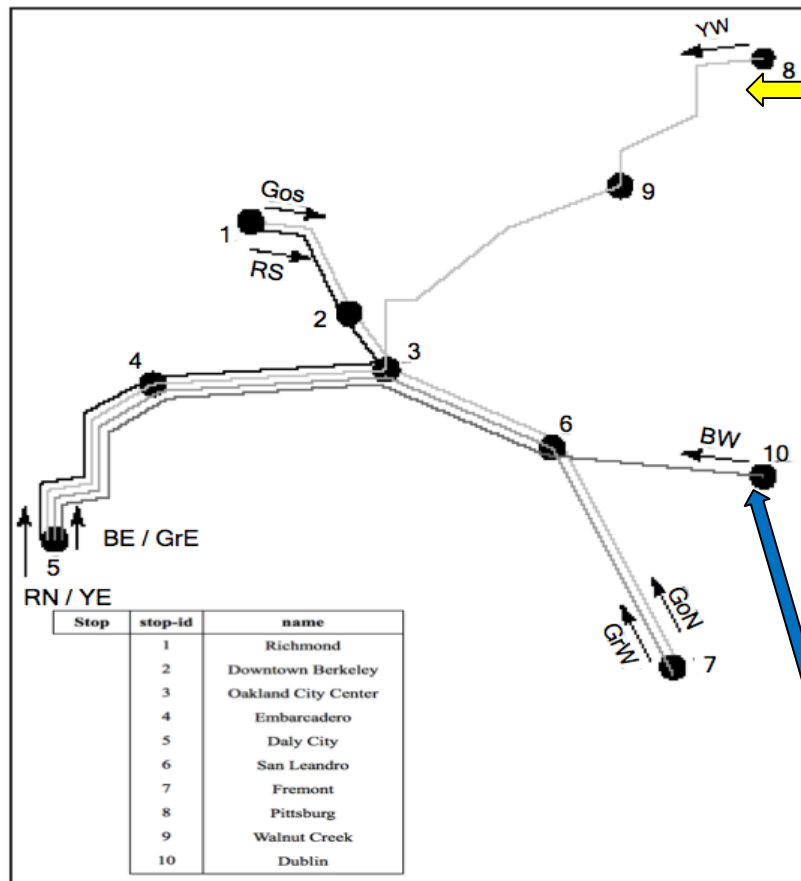
- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Spatial Network Query Example

- Find **shortest path** from a start-point to a destination
- Find **nearest** hospital by driving distance
- Find **shortest route** to deliver packages to a set of homes
- Allocate customers to **nearest** service center



Railway Network & Queries



- Find the number of stops on the Yellow West (YW) route
- List all stops which can be reached from Downtown Berkeley (2)
- List the routes numbers that connect Downtown Berkeley (2) & Daly City (5)
- Find the last stop on the Blue West (BW) route

River Network & Queries


- List the names of all direct and indirect tributaries (支流) of Mississippi river
- List the direct tributaries of Colorado
- Which rivers could be affected if there is a spill in North Platte river



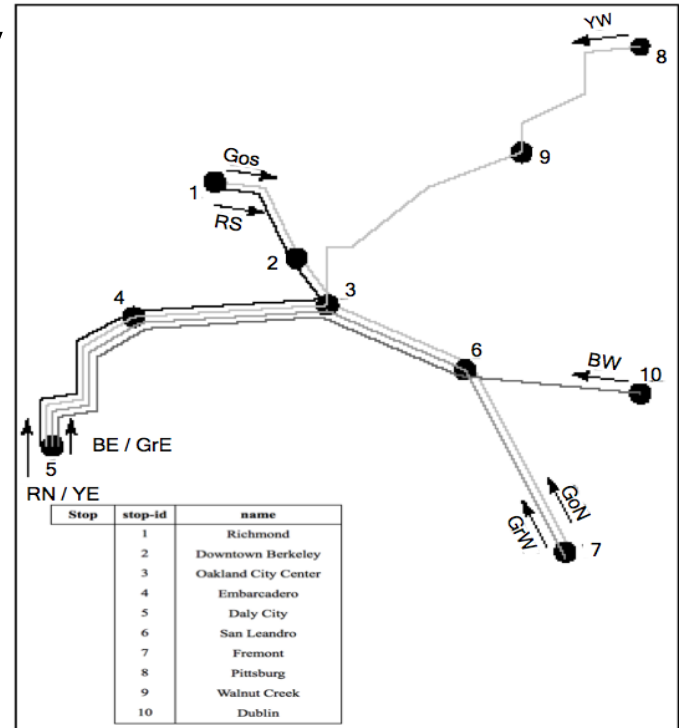
(b) River Network

Spatial Networks: Three Examples

A Road Network



A Railway Network



A River Network



(b) River Network

Quiz 3

- Which of the following are queries on a road network?
 - a) What is the shortest path to the nearest airport?
 - b) What is the driving distance to the nearest gas station?
 - c) How many road intersections are there in my city?
 - d) All of the above

第十一章 空间网络模型与查询

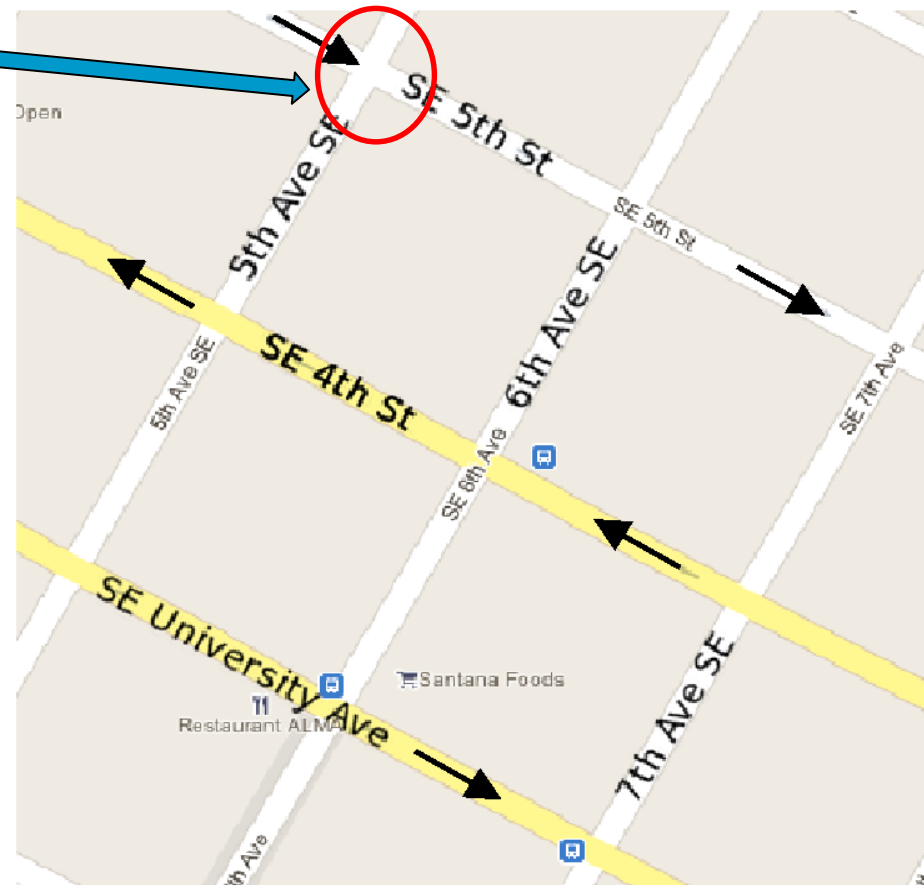
- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Data Models of Spatial Networks


- **Conceptual Model**
 - Information Model: Entity Relationship Diagrams
 - Mathematical Model: Graphs
- Logical Data Model
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage-Structures, File-Structures
 - Algorithms for common operations

Modeling Roadmaps

- Many Concepts, e.g.
 - Roads (or streets, avenues)
 - Road-Intersections
 - Road-Segments
 - Turns
 - ...





An Entity Relationship Diagram



Graph Models

- A Simple Mathematical Model
 - A graph $G = (V, E)$
 - V = a finite set of vertices
 - E = a set of edges model a binary relationship between vertices
- Example





A Graph Model of River Network

- Nodes = rivers
- Edges = a river falls into another river



(b) River Network





Pros and Cons of Graph Models

- Strength
 - Well developed mathematics for reasoning
 - Rich set of computational algorithms and data-structures
- Weakness
 - Models only one binary relationship
- Implications
 - A. Difficult to model multiple relationships, e.g., connect, turn
 - B. Multiple graph models possible for a spatial network



Modeling Turns in Roadmaps

- Approach 1: Model turns as a set of connects



- Approach 2: Use hyper-edges (and hyper-graphs)
- Approach 3: Annotate graph node with turn information

Alternative Graph Models for Roadmaps

- Choice 1:
 - Nodes = road-intersections
 - Edge (A, B) = road-segment **connects** adjacent road-intersections A, B
- Choice 2:
 - Nodes = (directed) road-segments
 - Edge (A, B) = **turn** from road-segment A to road-segment B
- Choice 3:
 - Nodes = roads
 - Edge(A, B) = road A **intersects_with** road B

Quiz 4

- Which of the following are usually not captured in common graph models of roadmaps?
 - a) Turn restrictions (e.g., no left turn)
 - b) Road intersections
 - c) Road segments
 - d) All of the above

第十一章 空间网络模型与查询

- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Data Models of Spatial Networks


- Conceptual Model
 - Information Model: Entity Relationship Diagrams
 - Mathematical Model: Graphs
- Logical Data Model & Query Languages
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage-Structures, File-Structures
 - Algorithms for common operations

Transitive Closure

- Consider a graph $G = (V, E)$
- Transitive closure(G) = $G^* = (V^*, E^*)$, where
 - $V^* = V$
 - (A, B) in E^* if and only if there is a path from A to B in G

Transitive Closure - Example

- Example
 - G has 5 nodes and 5 edges
 - G^* has 5 nodes and 9 edges
 - Note edge (1,4) in G^* for
 - path (1, 2, 3, 4) in G




(a) Graph G

R

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3

(b) Relation form



(c) Transitive closure (G) = Graph G

X

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3
1	3
2	4
5	4
1	4

(d) Transitive closure in relation form

Limitations of Original SQL


- Recall Relation algebra based languages
 - Ex. Original SQL
 - Can not compute transitive closure, e.g., shortest path

Supporting Graphs in SQL

- Abstract Data Type (user defined)
 - SQL3
 - May include shortest path operation!
- Custom Statements
 - SQL2 - CONNECT clause in SELECT statement
 - For **directed acyclic graphs**, e.g. hierarchies
 - SQL3 - WITH RECURSIVE statement
 - Transitive closure on general graphs
 - SQL3 - User defined data types
 - Can include shortest path operation!

Quiz 5

- Which of the following is not in the transitive closure of the following graph?
 - a) (1, 5)
 - b) (1, 4)
 - c) (2, 3)
 - d) (3, 4)



第十一章 空间网络模型与查询

- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Querying Graphs: Overview

- Relational Algebra
 - Can not express transitive closure queries
- Two ways to extend SQL to support graphs
 - Abstract Data Types
 - Custom Statements
 - SQL2 - CONNECT BY clause(s) in SELECT statement
 - SQL3 - WITH RECURSIVE statement
 - SQL3 - User defined data types


CONNECT BY: Input, Output

- Input:

- (a) Edges of a directed acyclic graph G
- (b) Start Node S, e.g., Missouri
- (c) Travel Direction

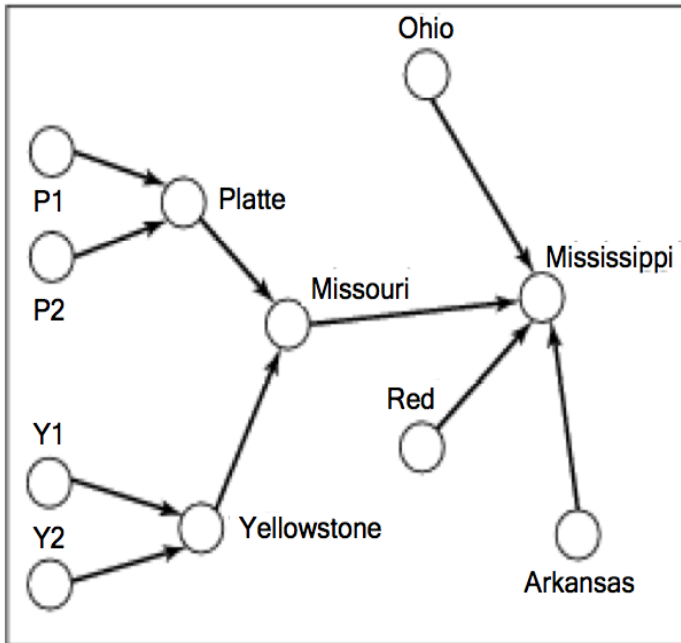
- Output: Transitive closure of G

- Ex. Predecessors of S = Missouri
- Ex. Successors of S = Missouri
(密苏里)



(a) Mississippi network (Y1 = Bighorn river, Y2 = Power river, P1 = Sweet water River, P2 = Big Thompson river)

Directed Edges: Tabular Representation




(a) Mississippi network (Y1 = Bighorn river, Y2 = Power river, P1 = Sweet water River, P2 = Big Thompson river)

Table: Falls_Into

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

CONNECT BY– PRIOR - START WITH

```
SELECT source
FROM Falls_Into
CONNECT BY PRIOR source = dest
START WITH dest = "Missouri"
```




Q? What does CONNECT BY ...
PRIOR specify?

- Direction of travel
- Example: **From Dest to Source**
- Alternative: **From Source to Dest**

Table: Falls_Into

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi



CONNECT BY– PRIOR - START WITH

- Syntax details
 - FROM clause a table for directed edges of an acyclic graph
 - PRIOR identifies direction of traversal for the edge
 - START WITH specifies first vertex for path computations
- Semantics
 - List all nodes reachable from first vertex using directed edge in specified table
 - Assumption - no cycle in the graph!
 - Not suitable for train networks, road networks

CONNECT BY– PRIOR - START WITH

Choice 1: Travel from Dest to Source

Ex. List direct & indirect tributaries of Missouri.

```
SELECT source
FROM   Falls_Into
CONNECT BY PRIOR source = dest
START WITH dest = "Missouri"
```

Choice 2: Travel from Source to Dest

Ex. Which rivers are affected by spill in Missouri?

```
SELECT dest
FROM   Falls_Into
CONNECT BY source = PRIOR dest
START WITH source = "Missouri"
```

Table: Falls_Into

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

Execution Trace – Step 1

```
SELECT source
FROM   Falls_Into
CONNECT BY PRIOR source = dest
START WITH dest = "Missouri"
```

1. Prior Result = SELECT * FROM Falls_Into
WHERE dest = "Missouri"

Table: "Prior "	
Source	Dest
Platte	Missouri
Yellowstone	Missouri




Table: Falls_Into	
Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

Execution Trace – Step 2

```
SELECT source
FROM Falls_Into
CONNECT BY PRIOR source = dest
START WITH dest = "Missouri"
```

2. Iteratively add `Join(Prior_Result.source = Falls_Into.dest)`

```
Prior Result = SELECT * FROM Falls_Into
                WHERE Prior_Result.source = Falls_Into.dest
```

Prior Result

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri

Prior Result

Source	Dest
Platte	Missouri
Yellowstone	Missouri

Table: Falls_Into


Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

SQL CONNECT Exercise

- Study 2 SQL queries on right
 - Note different use of PRIOR keyword
- Compute results of each query
- Which one returns ancestors of 3?
- Which returns descendants of 3?
- Which query lists river affected by
 - oil spill in Missouri (id = 3)?

SELECT source **FROM** FallsInto
CONNECT BY PRIOR source = dest
START WITH dest = 3

SELECT source **FROM** FallsInto
CONNECT BY source = **PRIOR** dest
START WITH dest = 3



Quiz 6

- Which of the following is false about CONNECT BY clause?
 - a) It is only able to output predecessors, but not successors, of the start node
 - b) It is able to output transitive closure of a directed graph
 - c) It usually works with PRIOR and START WITH keywords
 - d) None of the above

第十一章 空间网络模型与查询

- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Querying Graphs: Overview

- Relational Algebra
 - Can not express transitive closure queries
- Two ways to extend SQL to support graphs
 - Abstract Data Types
 - Custom Statements
 - SQL2 - CONNECT BY clause(s) in SELECT statement
 - **SQL3 - WITH RECURSIVE statement**
 - SQL3 - User defined data types

WITH RECURSIVE: Input, Output


- Input:

- (a) Edges of a directed graph G
- (b) Sub-queries to
 - Initialize results
 - Recursively grow results
 - Additional constraints

R

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3

(b) Relation form



(a) Graph G

- Output: Transitive closure of G

- Ex. Predecessors of a node
- Ex. Successors of a node

Syntax of WITH RECURSIVE Statement

WITH RECURSIVE X(source, dest)

← Description of Result Table

AS (SELECT source, dest FROM R)

← Initialization Query


UNION

(SELECT R.source, X.dest
FROM R, X
WHERE R.dest = X.source)

← Recursive Query to grow result

Example Input and Output

WITH RECURSIVE X(source,dest)
AS (**SELECT** source,dest **FROM** R)
 UNION
 (**SELECT** R.source, X.dest
FROM R, X
WHERE R.dest=X.source)




(a) Graph G

R

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3

(b) Relation form




(c) Transitive closure (G) = Graph G

X

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3
1	3
2	4
5	4
1	4

(d) Transitive closure in relation form


SQL3 Recursion Example - Meaning



(a) Graph G

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3

(b) Relation form



(c) Transitive closure (G) = Graph G

SOURCE	DEST
1	2
1	5
2	3
3	4
5	3
1	3
2	4
5	4
5	4
1	4

(d) Transitive closure in relation form

- Initialize X by
(SELECT source,dest FROM R)
 - Recursively grow X by
(SELECT R.source, X.dest
FROM R, X
WHERE R.dest=X.source)
 - Infer X(a,c) from R(a,b),X(b,c)
- Infer X(1,3) from R(1,2),X(2,3)
 Infer X(2,4) from R(2,3),X(3,4)
 Infer X(5,4) from R(5,3),X(3,4)
 Infer X(1,4) from R(1,5),X(5,4)

SQL3 Recursion

- Syntax
 - WITH RECURSIVE <Relational Schema>
 - AS <Query to populate relational schema> Syntax details
 - <Relational Schema> lists columns in result table with directed edges
 - <Query to populate relational schema> has UNION of nested sub-queries
 - Base cases to initialize result table
 - Recursive cases to expand result table

SQL3 Recursion

- Semantics

- Results relational schema say $X(\text{source}, \text{dest})$
 - Columns source and dest come from same domain, e.g. Vertices
 - X is a edge table, $X(a,b)$ directed from a to b
- Result table X is initialized using base case queries
- Result expanded using $X(a, b)$ and $X(b, c)$ implies $X(a, c)$

With Recursive for Connectivity

- Connect by
 - For directed acyclic graphs, e.g. hierarchies
 - PostgreSQL, SQL Server not supported
 - Oracle supported
- With Recursive
 - Transitive closure on general graphs
 - PostgreSQL, SQL Server, Oracle supported

<http://www.postgresql.org/docs/current/static/queries-with.html>

Case Studies

- Goal: Compare relational schemas for spatial networks
 - River networks has an edge table, Falls_Into
 - BART train network does not an edge table
 - Edge table is crucial for using SQL transitive closure
 - Exercise: Proposed a different set of table to model BART as a graph
 - using an edge table connecting stops
- River networks - graph model
 - Can use SQL transitive closure to compute ancestors or descendent of a river
 - We saw an examples using CONNECT BY clause
 - Exercises explore use of WITH RECURSIVE statement

Case Studies

- BART train network - non-graph model
 - Entities = Stop, Route
 - Relationship = aMemberOf(Stop, Route)
 - Can not use SQL recursion
 - No table can be viewed as edge table
 - RouteStop table is a subset of transitive closure
- Transitive closure queries on edge(from_stop, to_stop)
 - A few can be answered by querying RouteStop table
 - Many can not be answered
 - Find all stops reachable from Downtown Berkeley

Quiz 7

- Which of the following are true about WITH RECURSIVE clause?
 - a) It is able to output transitive closure of a directed graph
 - b) It usually works with an edge table
 - c) It includes two SELECT statements
 - d) All of the above

第十一章 空间网络模型与查询

- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

With Recursive in PostgreSQL

- With clause
 - A way to write auxiliary statements for use in a larger query
 - Define temporary tables that exist just for one query
- A WITH query can refer to its own output
- Sum the integers from 1 through 100

```
WITH RECURSIVE t(n) AS (
```

```
    VALUES (1)
```

```
    UNION ALL
```

```
    SELECT n+1 FROM t WHERE n < 100
```

```
)
```

```
SELECT sum(n) FROM t;
```

A non-recursive term

Union or Union ALL

A recursive term

Recursive Query Evaluation

- Step 1: Evaluate the non-recursive term
 - For UNION, discard duplicate rows
 - Include all remaining rows in **the result of the recursive query**, and also place them in **a temporary working table**
 - Example
 - VALUES (1)

```
WITH RECURSIVE t(n) AS (  
  VALUES (1)  
  UNION ALL  
  SELECT n+1 FROM t WHERE n < 100)
```

Query Result	Working
1	1

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - Example
 - `SELECT n+1 FROM t WHERE n < 100`
 - t is the query result table Q?
or the working table W

```
WITH RECURSIVE t(n) AS (  
  VALUES (1)  
  UNION ALL  
  SELECT n+1 FROM t WHERE n < 100)
```

Intermediate
2

Query Q
1

Working W
1

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - 2.2 For UNION, discard duplicated rows and rows that duplicated any previous result row
 - Where are previous result rows?

```
WITH RECURSIVE t(n) AS (  
  VALUES (1)  
  UNION ALL  
  SELECT n+1 FROM t WHERE n < 100)
```

Intermediate	Query Q	Working W
2	1	1

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - 2.2 For UNION, discard duplicated rows and rows that duplicated any previous result row
 - 2.3 Including all remaining rows in **the result of the recursive query**, and also place them in **a temporary intermediate table**

Intermediate	Query Q	Working W
2	1	1
	2	

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - 2.2 For UNION, discard duplicated rows and rows that duplicated any previous result row
 - 2.3 Including all remaining rows in **the result of the recursive query**, and also place them in **a temporary intermediate table**
 - 2.4 Replace the contents of **the working table** with the contents of **the intermediate table**, then empty **the intermediate table**

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps
 - SELECT $n+1$ FROM **W** WHERE $n < 100$

Intermediate	Query Q	Working W
	1	2
	2	
Intermediate	Query Q	Working W
3	1	2
	2	
Intermediate	Query Q	Working W
	1	3
	2	
	3	

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps
 - SELECT $n+1$ FROM **W** WHERE $n < 100$

```
WITH RECURSIVE t(n) AS (  
  VALUES (1)  
  UNION ALL  
  SELECT  $n+1$  FROM t WHERE  $n < 100$ )
```

t
1
.....
100

Intermediate	Query Q	Working W
	1	100
	
	100	

Intermediate	Query Q	Working W
	1	
	
	100	

Recursive Query Evaluation in C

- with recursive $X(\dots)$ as (
 - $SQL_A(O)$
 - $union / union\ all$
 - $SQL_B(O, X)$)
- $X(\dots) = SQL_A[remove\ duplicates\ for\ union];$
 $W(\dots) = X(\dots);$
while (table W is not empty) {
 - $T(\dots) = SQL_B(O, W) [except\ X(\dots)\ for\ union];$
 - $X(\dots) = X(\dots) union / union\ all\ T(\dots);$
 - $W(\dots) = T(\dots);$}

O: 原始关系/表
X: 最终关系/表
W: 临时关系/表
T: 临时关系/表

问题: Union与Union ALL区别
避免使用union all

Recursive Query Evaluation in C

WITH RECURSIVE $t(n)$ AS (

VALUES (1)

UNION ALL

SELECT $n+1$ FROM t WHERE $n < 100$)

t : 最终关系/表

W : 临时关系/表

T : 临时关系/表

- insert into $t(n)$ values(1);

delete from W ; insert into $W(n)$ select * from t ;

while ((select count(*) from W) \neq 0) {

delete from T ;

insert into $T(n)$ select $n+1$ from W where $n < 100$;

insert into $t(n)$ select * from T ; // union all

delete from W ; insert into W from select * from T ;

}

With Recursive Limitation


- Strictly speaking, this process is **iteration** not recursion, but RECURSIVE is the terminology chosen by the SQL standards committee
- Recursive queries are typically used to deal with **hierarchical** or **tree-structured data**

With Recursive Limitation

- Important: the recursive part of the query will eventually return no tuples, or else the query will **loop indefinitely**
 - Using UNION instead of UNION ALL can accomplish this by discarding rows that duplicate previous output rows
- A cycle does not involve output rows that are completely duplicate
 - It may be necessary to check just one or a few fields to see if the same point has been reached before
- Standard method
 - Compute **an array** of the already-visited values


With Recursive Example 1

- 广度优先遍历 - 深度depth
- 数据库中，右图的关系为edges(start, end)
 - 6行记录(A,B), (A,C), (B,A), (B,C), (C,A), (C,B)
- with recursive X(node, depth) as (
 select start, 0 from edges where start = A
union
 select end, depth + 1 from edges, X
 where start = node and depth < 3)
- 如果不加depth<3，上述语句会死循环



With Recursive Example 1

- with recursive $X(\text{node}, \text{depth})$ as (
select start, 0 from edges where start = A
union
select end, depth + 1 from edges, X
where start = node and depth < 3)




- 初始X: (A, 0)
- 第1次迭代: (A, 0), (B, 1), (C, 1)
- 第2次迭代: (A, 0), (B, 1), (C, 1), (A, 2), (B, 2), (C, 2)
 - 如果使用union all得到: (A, 2), (C, 2), (A, 2), (B, 2)
- 第3次迭代: (A, 0), (B, 1), (C, 1), (A, 2), (B, 2), (C, 2), (A, 3), (B, 3), (C, 3)
-

With Recursive Example 2


- 避免重复遍历节点path和cycle
- with recursive X(node, depth, path, circle) as (
 select start, 0, array[start], false from edges where start = A
union
 select end, depth + 1, path || end, end = any(path)
from edges, X
where start = node and not circle)
- PostgreSQL中的数组
 - 初始化: array[...]
 - 增加元素: path || end
 - 判断元素是否在数组中: end = any(path)

With Recursive Example 2

- with recursive X(node, depth, path, circle) as (
select start, 0, array[start], false from edges where start = A
union
select end, depth + 1, path || end, end = any(path)
from edges, X
where start = node and not circle)
- 初始X: (A, 0, [A], false)
- 第1次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false)
- 第2次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false)
- 第3次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false),
(A, 3, [A, B, C, A], true), (B, 3, [A, B, C, B], true),
(A, 3, [A, C, B, A], true), (C, 3, [A, C, B, C], true)



With Recursive Example 2



- 初始X: (A, 0, [A], false)
- 第1次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false)
- 第2次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false)
- 第3次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false),
(A, 3, [A, B, C, A], true), (B, 3, [A, B, C, B], true),
(A, 3, [A, C, B, A], true), (C, 3, [A, C, B, C], true)
- 第4次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false),
(A, 3, [A, B, C, A], true), (B, 3, [A, B, C, B], true),
(A, 3, [A, C, B, A], true), (C, 3, [A, C, B, C], true)

● 结束

With Recursive Example 3

- Find all the direct and indirect sub-parts of a product, given only a table that shows immediate inclusions
 - parts(part, sub_part, quantity)
 - Similar to Rivers

With Recursive Example 3

- Find all the direct and indirect sub-parts of a product

```
WITH RECURSIVE included_parts(sub_part, part, quantity) AS (  
  SELECT sub_part, part, quantity FROM parts WHERE part = 'our_product'  
  UNION ALL  
  SELECT p.sub_part, p.part, p.quantity FROM included_parts pr, parts p  
  WHERE p.part = pr.sub_part  
)
```

```
SELECT sub_part, SUM(quantity) as total_quantity  
FROM included_parts  
GROUP BY sub_part
```

With Recursive Example 4

- Graph search: graph(id, link, data)

```
WITH RECURSIVE search_graph(id, link, data, depth) AS (  
    SELECT g.id, g.link, g.data, 1  
    FROM graph g  
    UNION ALL  
    SELECT sg.id, g.link, g.data, sg.depth + 1  
    FROM graph g, search_graph sg  
    WHERE g.id = sg.link  
)  
SELECT * FROM search_graph;
```

- This query will loop if the link relationship contain cycle
 - UNION ALL → UNION
 - Would not eliminate the looping due to the “depth” output

With Recursive Example 4

- Need to recognize whether we have reached the same row again while following path of links
 - Add two columns path and cycle to the loop-prone query

WITH RECURSIVE search_graph(id, link, data, depth, path, cycle) **AS** (

SELECT g.id, g.link, g.data, 1,

ARRAY[g.id],

false

FROM graph g

UNION ALL

SELECT sg.id, g.link, g.data, sg.depth + 1,

path || g.id,

g.id = **ANY**(path)

FROM graph g, search_graph sg

WHERE g.id = sg.link **AND** NOT cycle

Represent the “path” taken to reach any particular row

)

With Recursive Example 4

- Generally, more than one field needs to be checked to recognize a cycle, use an array of rows

```
WITH RECURSIVE search_graph(id, link, data, depth, path, cycle) AS (  
    SELECT g.id, g.link, g.data, 1,  
           ARRAY[ROW(g.f1, g.f2)],  
           false  
    FROM graph g  
    UNION ALL  
    SELECT g.id, g.link, g.data, sg.depth + 1,  
           path || ROW(g.f1, g.f2),  
           ROW(g.f1, g.f2) = ANY(path)  
    FROM graph g, search_graph sg  
    WHERE g.id = sg.link AND NOT cycle  
)  
SELECT * FROM search_graph;
```

Recursive in PostgreSQL

- A helpful trick for testing queries when you are not certain if they might loop is to place a **LIMIT** in the parent query

```
WITH RECURSIVE t(n) AS (  
    VALUES (1)  
    UNION ALL  
    SELECT n+1 FROM t WHERE n < 100  
)  
SELECT sum(n) FROM t LIMIT 100;
```

- PostgreSQL evaluates only as many rows of a WITH query as are actually fetched by the parent query
 - Other DBMS might work different
 - Won't work if order by or join them to some other table

With Queries Applications

- Avoid redundant work
 - Evaluated only once per execution of the parent query, even if they are referred to more than once by the parent query or sibling WITH queries
 - Expensive calculations that are needed in multiple places can be placed within a WITH query
- Prevent unwanted multiple evaluations of functions with side-effects
 - The optimizer is less able to push restrictions from the parent query down into a WITH query than an ordinary sub-query

第十一章 空间网络模型与查询


- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Data Models of Spatial Networks

- Conceptual Model
 - Entity Relationship Diagrams, Graphs
- Logical Data Model & Query Languages
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage: Data-Structures, File-Structures
 - Algorithms for common operations


Main Memory Data-Structures

- Adjacency matrix
 - $M[A, B] = 1$ if and only if $\text{edge}(\text{vertex } A, \text{vertex } B)$ exists
- Adjacency list :
 - Maps a vertex to a list of its successors



Disk-based Tables

- Normalized tables
 - One for vertices, other for edges
- Denormalized
 - One table for nodes with adjacency lists



Spatial Network Storage

- Problem Statement
 - Given a spatial network
 - Find efficient data-structure to store it on disk sectors
 - Goal - Minimize I/O-cost of operations
 - Find(), Insert(), Delete(), Create()
 - Get-A-Successor(), Get-Successors()
 - Constraints
 - Spatial networks are much larger than main memories
- Problems with Geometric indices, e.g. R-tree
 - Clusters objects by proximity not edge connectivity
 - Performs poorly if edge connectivity not correlated with proximity (邻近)
- Trends: graph based methods


Graph Based Storage Methods

- Insight:
 - I/O cost of operations (e.g. get-a-successor) minimized by maximizing CRR
 - $CRR = \text{Pr. (node-pairs connected by an edge are together in a disk sector)}$

File-Structures:


Partition Graph into Disk Blocks

- Which partitioning reduces disk I/O for graph operations?
 - Choice 1: Geometric partition
 - Choice 2: Min-cut Graph Partition
 - Choice 2 cuts fewer edges and is preferred
 - Assuming uniform querying popularity across edges




Sample Network

○ Node
— Edge



Sample Network
Different data pages in different colors

— Cut edge


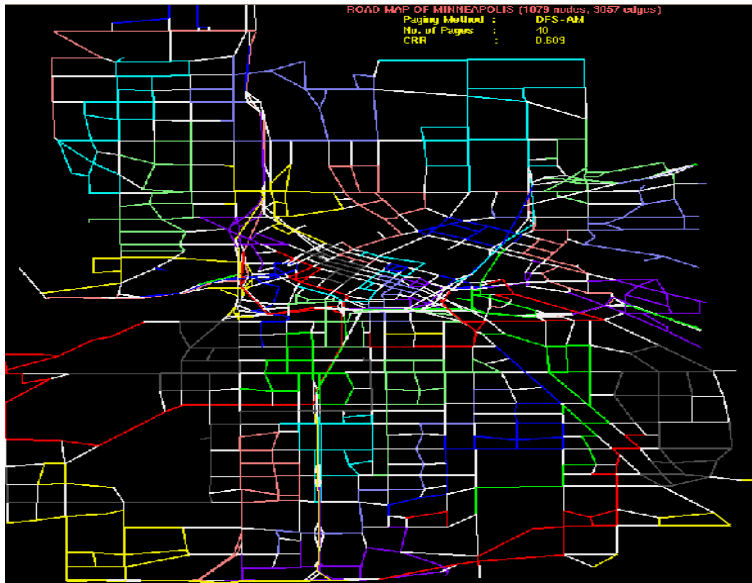


Sample Network
Different data pages in different colors

— Cut edge

Graph Based Storage Methods

- Consider two disk-paging of Minneapolis (明尼阿波利斯) major roads
 - Non-white edges \Rightarrow node pair in same page
 - White edge are cut-edges
 - Node partitions on right has fewer cut-edges and is preferred \Rightarrow higher CRR





Clustering and Storing a Sample Network

- Storage method idea
 - Divide nodes into sectors
 - to maximize CRR
 - Use a secondary index
 - for find()
 - using R-tree or B-tree


- Example

- left part : node division
- right part
 - disk sectors
 - secondary index
 - B-tree/Z-order



Exercise: Graph Based Storage Methods

- If a disk page holds 3 records, which partitioning will has fewest cut-edges?
 - (1, 2, 3), (4, 5, 6)
 - (2, 3, 4), (1, 5, 6)
 - (1, 2, 6), (3, 4, 5)
 - (1, 3, 5), (2, 4, 6)



Node

nid	x	y	Successors	Predecessors
1	—	—	(2,5,6)	()
2	—	—	(3,5)	(1)
3	—	—	(4)	(3)
4	—	—	(5)	(3)
5	—	—	(6)	(2,1)
6	—	—	()	(1,5)

Quiz 8

- Which of the following is not disk-based representations of graphs?
 - a) Normalized tables (e.g., node table and edge table)
 - b) Denormalized table (e.g., node table with successors and predecessors columns)
 - c) Adjacency matrix
 - d) All of the above

第十一章 空间网络模型与查询

- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Data Models of Spatial Networks

- Conceptual Model
 - Entity Relationship Diagrams, Graphs
- Logical Data Model & Query Languages
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage: Data-Structures, File-Structures
 - Algorithms for common operations

Query Processing for Spatial Networks

- Query Processing
 - DBMS decomposes a query into building blocks
 - Keeps a couple of strategy for each building block
 - Selects most suitable one for a given situation
- Building blocks
 - Connectivity(A, B): Is node B reachable from node A?
 - Shortest path(A, B): Identify the least cost path from node A to node B

Algorithms

- Main memory
 - Connectivity: Breadth first search, Depth first search
 - Shortest path: Dijkstra's algorithm, A*
- Disk-based
 - Shortest path - Hierarchical routing algorithm
 - Connectivity strategies are in SQL3


Algorithms for Connectivity Query

- Breadth first search

- Visit descendent by generation
- Children before grandchildren
- Example: 1 - (2,4) - (3, 5)

- Depth first search

- Try a path till dead-end
- Backtrack to try different paths
- Like a maze game
- Example: 1-2-3-2-4-5
- Note backtrack from 3 to 2



Quiz 9

- Which of the following is false?
 - a) Breadth first search visits nodes layer (i.e. generation) by layer
 - b) Depth first search try a path till dead-end, then backtrack to try different paths
 - c) Depth first search always performs better than breadth first search
 - d) None of the above

第十一章 空间网络模型与查询


- 11.1 Motivation, and use cases
- 11.2 Example spatial networks
- 11.3 Conceptual model
- 11.4 Need for SQL extensions
- 11.5 CONNECT statement
- 11.6 RECURSIVE statement
- 11.7 RECURSIVE in PostgreSQL
- 11.8 Storage and data structures
- 11.9 Algorithms for connectivity query
- 11.10 Algorithms for shortest path

Shortest Path Algorithms

- Iterate
 - Expand most promising descent node
 - Dijkstra's: try closest descendent to self
 - A^* : try closest descendent to both destination and self
 - Update current best path to each node, if a better path is found
- Till destination node is expanded

Dijkstra's algorithm

- Dijkstra's algorithm
 - Identify paths to descendent by depth first search
 - Each iteration
 - Expand descendent with smallest cost path so far
 - Update current best path to each node, if a better path is found
 - Till destination node is expanded



Node (R)

id	x	y
1	4.0	5.0
2	6.0	3.0
3	5.0	1.0
4	3.0	2.0
5	1.0	3.0

Edge (S)

source	dest	distance
1	2	$\sqrt{8}$
1	4	$\sqrt{10}$
2	3	$\sqrt{5}$
2	4	$\sqrt{10}$
4	5	$\sqrt{5}$
5	1	$\sqrt{18}$

Dijkstra's algorithm

Node (R)


id	x	y
1	4.0	5.0
2	6.0	3.0
3	5.0	1.0
4	3.0	2.0
5	1.0	3.0

Edge (S)

source	dest	distance
1	2	$\sqrt{8}$
1	4	$\sqrt{10}$
2	3	$\sqrt{5}$
2	4	$\sqrt{10}$
4	5	$\sqrt{5}$
5	1	$\sqrt{18}$

- Example:

- Consider `shortest_path(1,5)` for graph
- Iteration 1 expands node 1 and edges (1,2), (1,4)
 - set $\text{cost}(1,2) = \sqrt{8}$; $\text{cost}(1,4) = \sqrt{10}$ using Edge table
- Iteration 2 expands least cost node 2 and edges (2,3), (2,4)
 - set $\text{cost}(1,3) = \sqrt{8} + \sqrt{5}$
- Iteration 3 expands least cost node 4 and edges (4,5)
 - set $\text{cost}(1,5) = \sqrt{10} + \sqrt{5}$
- Iteration 4 expands node 3 and Iteration 5 stops node 5
- Answer is the path (1-4-5)




A*

- Best first algorithm
 - Similar to Dijkstra's algorithm with one change
 - $\text{Cost}(\text{node}) = \text{actual_cost}(\text{source}, \text{node}) + \text{estimated_cost}(\text{node}, \text{destination})$
 - `estimated_cost` should be an underestimate of actual cost
 - Example - euclidean distance
- Given effective `estimated_cost()` function, it is faster than Dijkstra's algorithm

● Example:

- Revisit `shortest_path(1,5)` for graph in
- Iteration 1 expands node 1 and edges (1,2), (1,4)
 - set `actual_cost(1,2) = sqrt(8)`; `actual_cost(1,4) = sqrt(10)`;
 - `estimated_cost(2,5) = 5`; `estimated_cost(4,5) = sqrt(5)`
- Iteration 2 expands least cost node 4 and edges (4,5)
 - set `actual_cost(1,5) = sqrt(10) + sqrt(5)`, `estimated_cost(5,5) = 0`
- Iteration 3 expands node 5
- Answer is the path (1-4-5)



Node (R)


id	x	y
1	4.0	5.0
2	6.0	3.0
3	5.0	1.0
4	3.0	2.0
5	1.0	3.0

Edge (S)


source	dest	distance
1	2	$\sqrt{8}$
1	4	$\sqrt{10}$
2	3	$\sqrt{5}$
2	4	$\sqrt{10}$
4	5	$\sqrt{5}$
5	1	$\sqrt{18}$

Dijkstra's vs. A*

Wave 0:



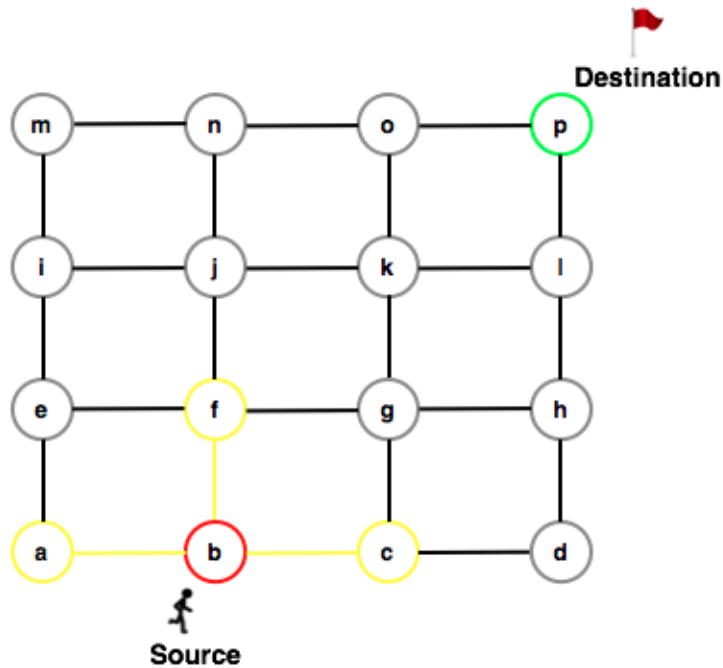
Dijkstra's Algorithm



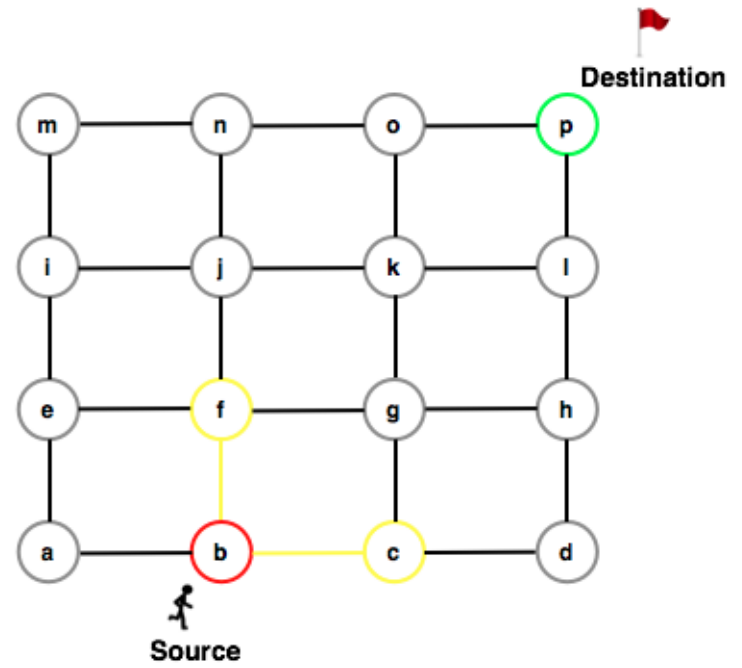
A* Algorithm

Dijkstra's vs. A*

Wave 1:



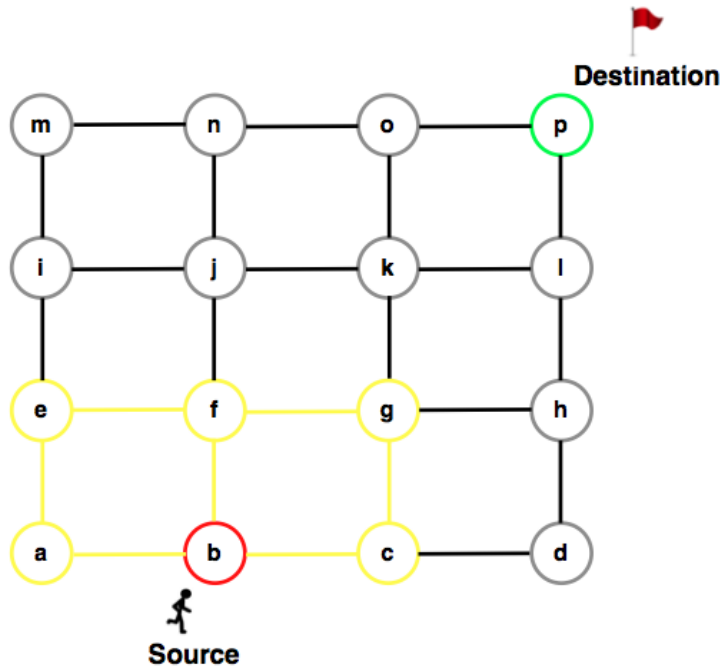
Dijkstra's Algorithm



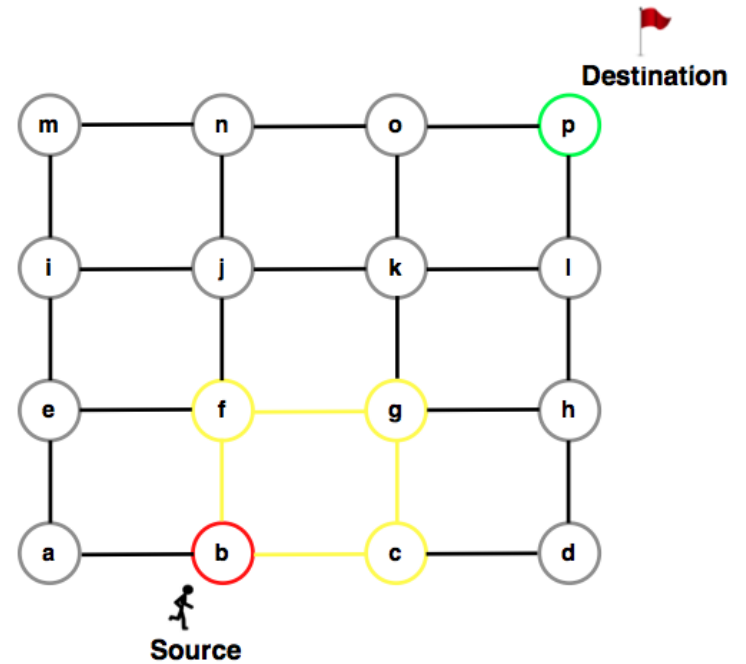
A* Algorithm

Dijkstra's vs. A*

Wave 2:



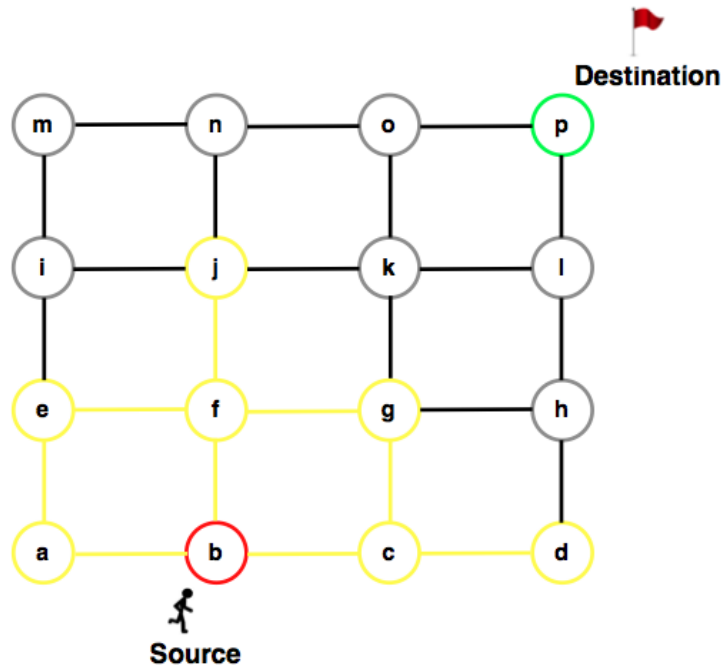
Dijkstra's Algorithm



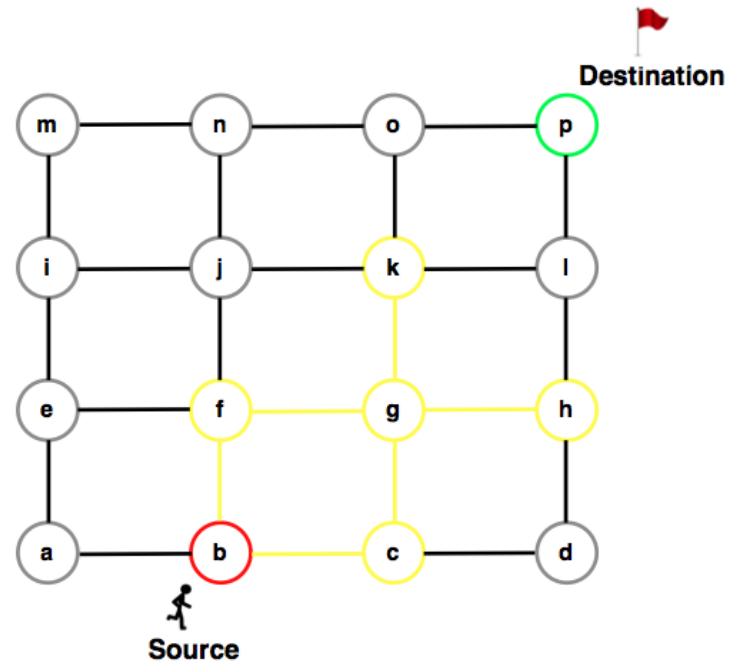
A* Algorithm

Dijkstra's vs. A*

Wave 3:



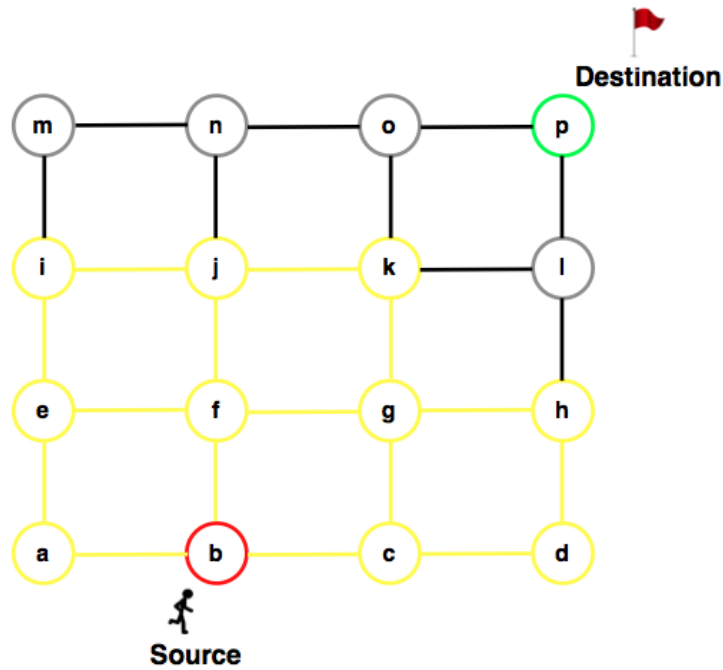
Dijkstra's Algorithm



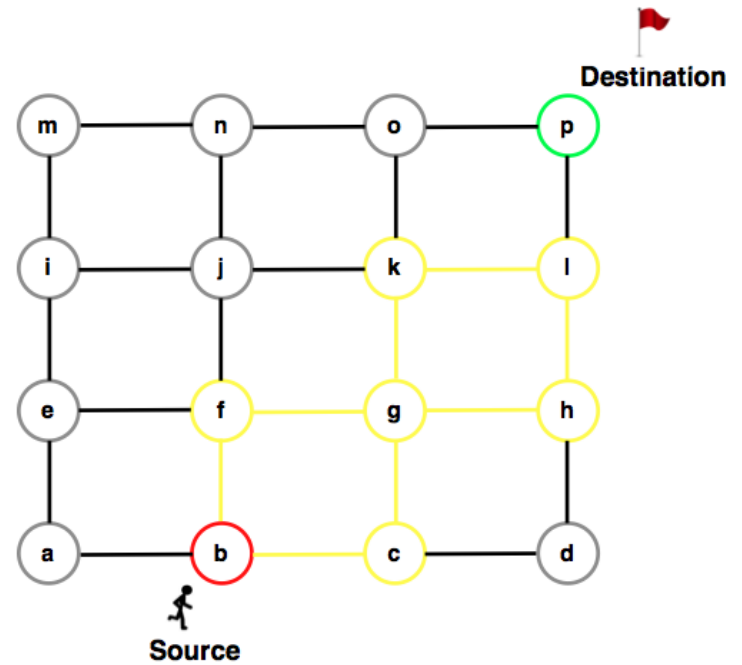
A* Algorithm

Dijkstra's vs. A*

Wave 4:



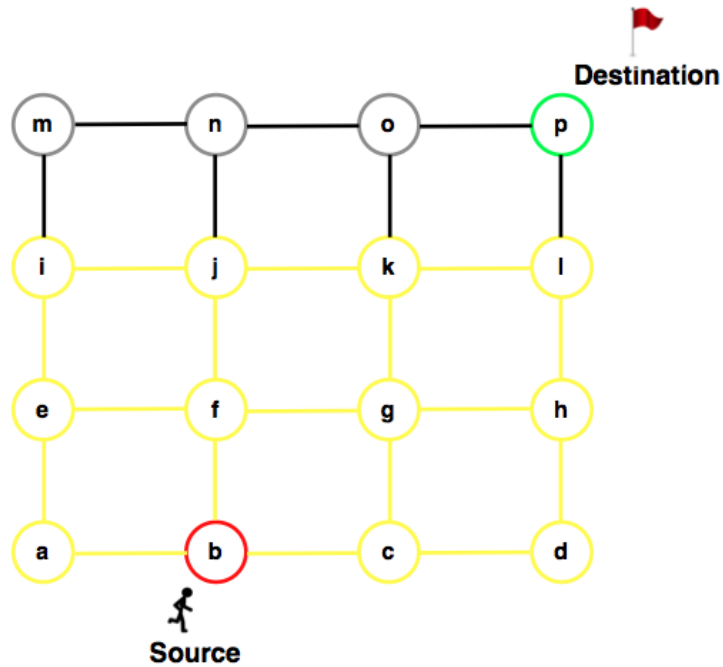
Dijkstra's Algorithm



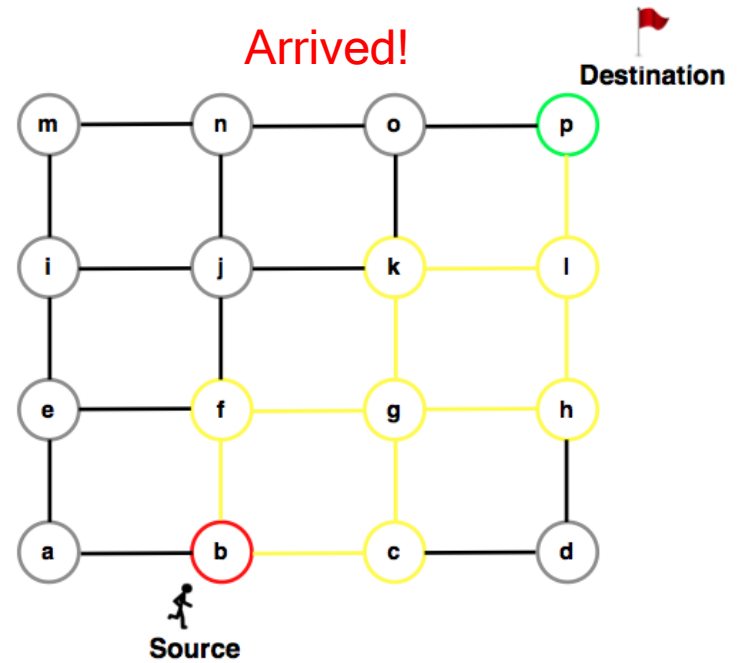
A* Algorithm

Dijkstra's vs. A*

Wave 5:

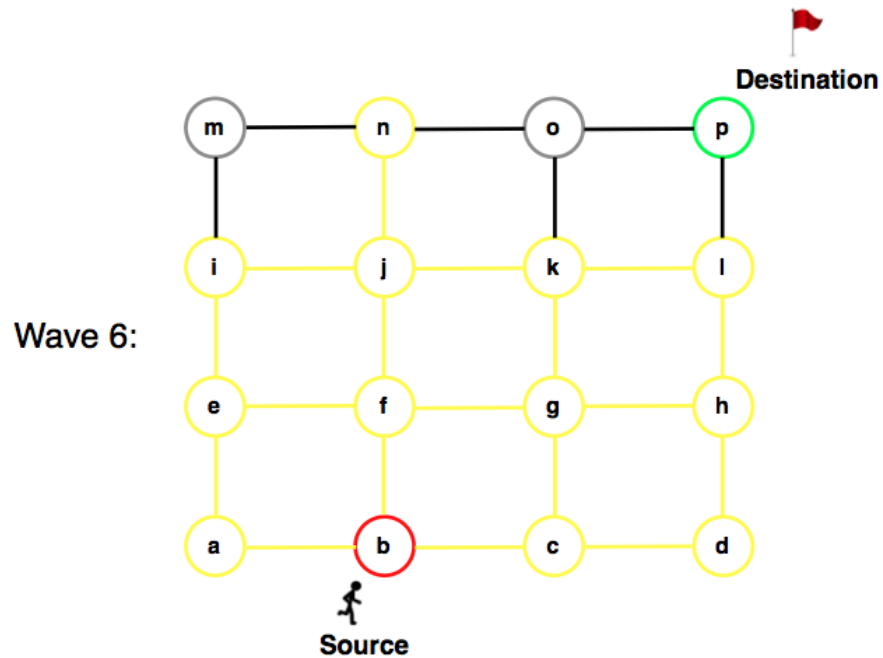


Dijkstra's Algorithm



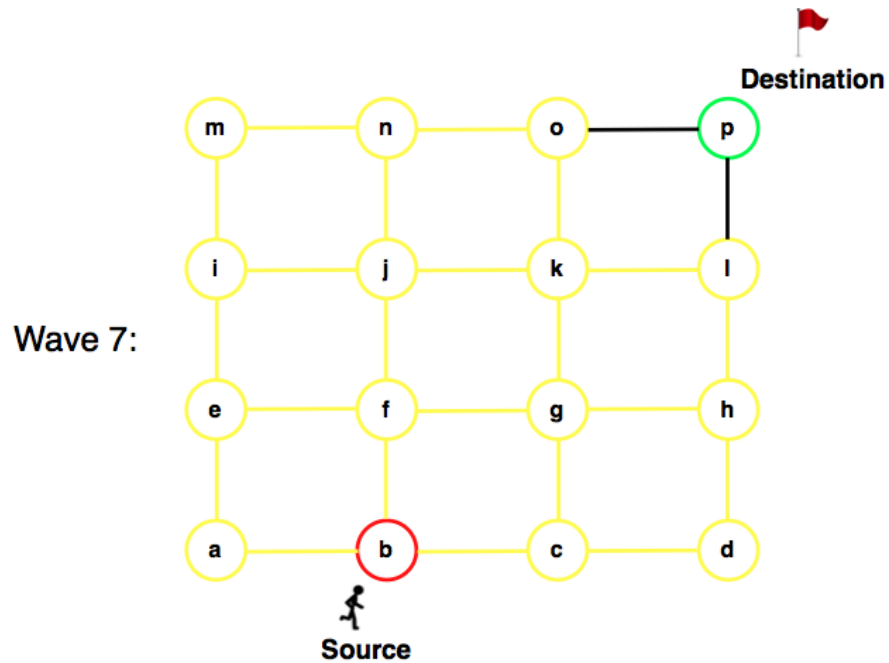
A* Algorithm

Dijkstra's vs. A*



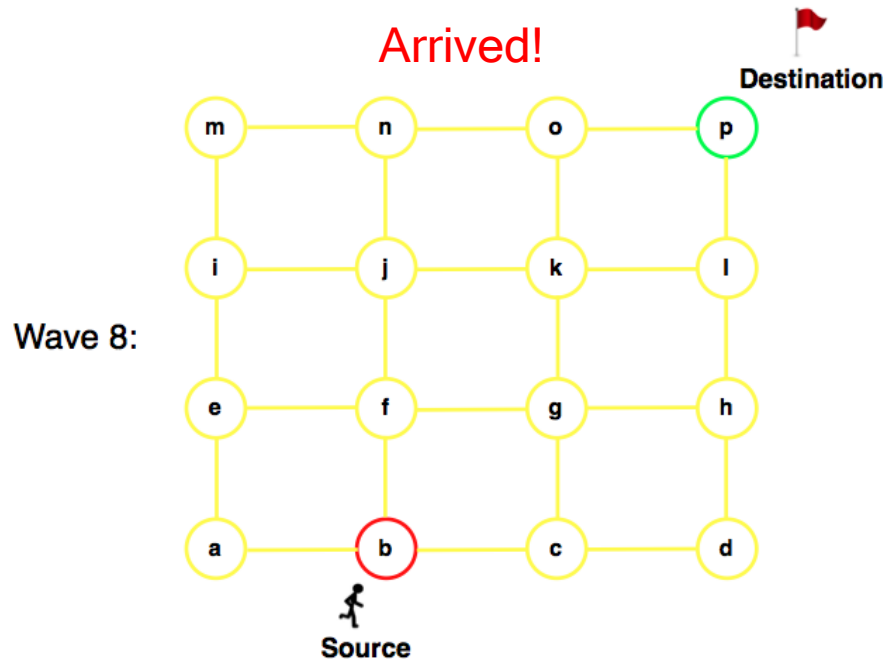
Dijkstra's Algorithm

Dijkstra's vs. A*



Dijkstra's Algorithm

Dijkstra's vs. A*



Dijkstra's Algorithm

Shortest Path Algorithms


- Iterate
 - Expand most promising node
 - Dijkstra's: try closest descendent to self
 - A^* : try closest descendent to both destination and self
 - Update current best path to each node, if a better path is found
- Till destination node is expanded
- Correct assuming
 - Sub-path optimality
 - Fixed, positive and additive edge costs
 - A^* : underestimate function

Shortest Path Strategies - 3

- Dijkstra's and Best first algorithms
 - Work well when entire graph is loaded in main memory
 - Otherwise their performance degrades substantially
- Hierarchical Routing Algorithms
 - Works with graphs on secondary storage
 - Loads small pieces of the graph in main memories
 - Can compute least cost routes

Hierarchical Routing: Simple Example


- Goal: Find the shortest path between Nodes 2 and 12
- Candidate algorithms: Dijkstra's algorithm, A*, and hierarchical routing



The edge length is 1 for every edge, except $\text{length}(\text{edge } (11-18)) = 3$

Trace Dijkstra's algorithm

- Tie-braking: prefer node with higher index



Q? Will A* (with Euclidean distance) be efficient?

18 nodes expanded


Expanded:	Open List:
2	1, 3, 5
5	1, 3, 6
3	1, 4, 6
1	4, 6
6	4, 7, 9
4	7, 9
9	7, 8
7	8, 10
10	8, 17
8	17
17	16, 18
18	15, 16, 19
16	13, 15, 19
19	13, 15
15	12, 13, 14
14	12, 13
13	11, 12
12	

Trace for A* algorithm

Cost(node n) = graph_distance (2, n) + Euclidean_distance (n, 12)

14 nodes expanded

Tie-breaking: prefer node with higher index




Q? How may we reduce computation cost further?

Expand:	Open List:
2	1, 3, 5
3	1, 4, 5
4	1, 5
5	1, 6
6	1, 7, 9
7	1, 9, 10
10	1, 9, 17
1	9, 17
9	8, 17
17	8, 16, 18
8	16, 18
18	15, 16, 19
15	12, 14, 16, 19
12	

Core Idea of the Hierarchical Routing

- Recognize **Islands** and **bridges**
- $SP(2, 12)$ must include **bridge** edge (10, 17)
- Divide n conquer : $SP(2, 12) = SP(2, 10) + \text{edge}(10, 17) + SP(17, 12)$
- Generalize to the case of multiple **bridges**



Trace Hierarchical Routing

4 nodes expanded in SP(17, 12) using A*

Expanded:

17
18
15
12

Open List:

16,18
15,16,19
12,14,16,19

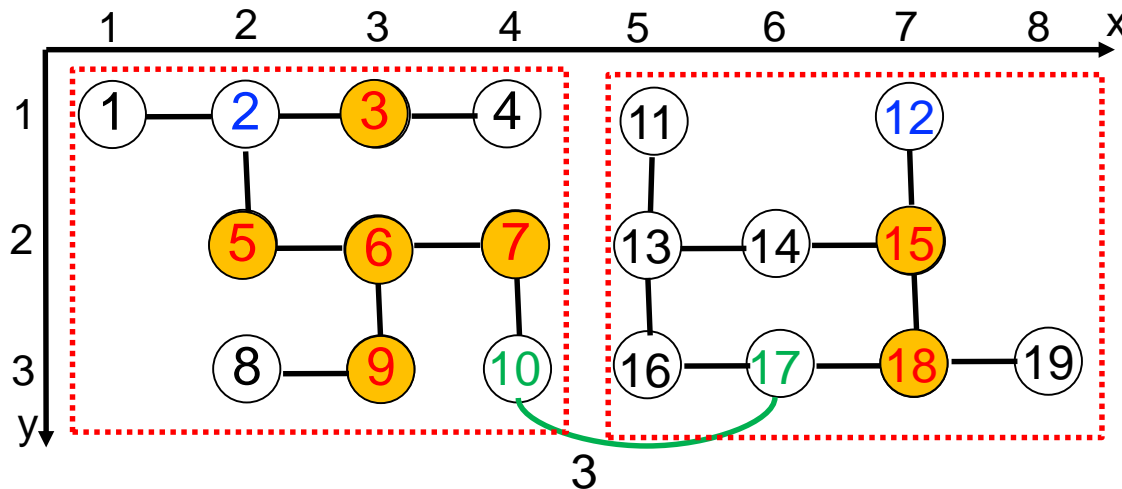
#Nodes expanded: $7 + 4 = 11$
7 nodes expanded in SP(2, 10) using A*

Expanded:

2
5
3
6
9
7
10


Open List:

1,3,5
1,3,6
1,4,6
1,4,7,9
1,4,7,8
1,4,8,10



Did We Reduce Computational Cost?


Algorithm (World View)	# of nodes expanded)
Dijkstra's (Graphs)	18
A* (Spatial Graphs)	14
Hierarchical Routing (Islands)	$7 + 4 = 11$



Q? What if Multiple bridges?
Q? How to choose among bridges?

Challenges: Multiple Islands & Bridges



- Invariant: $SP(2,12)$ must use one or more bridges
- **Challenge 1**: Multiple islands increase computational cost
- **Challenge 2**: Multiple bridges per island increase computational cost



Hierarchical Algorithm with Multiple Islands



Data Structures:

- Port node (a.k.a., boundary node) : a Node with edges to multiple islands
- Port Graph (a.k.a., boundary graph) :
- Island graphs (a.k.a., fragment graphs)
- Precompute & store
 - Shortest_path_costs: $SPC(\text{node } i, \text{node } j)$ for all (or selected) node pairs (i, j)



Hierarchical Algorithm with Multiple Islands

- Choose port pair (a.k.a., boundary node pair) for $SP(2,12)$
 - $\min SPC(2, \text{local port}) + SPC(2\text{'s local port}, 12\text{'s local port}) + SPC(\text{local port}, 12)$
 - Choices: (2,11), (2, 17), (4, 11), (4, 17)
 - Chosen port pair: (4, 11)
- Divide and conquer: $SP(2,12) = SP(2, 4) \cdot SP(4, 11) \cdot SP(11, 12)$
 - Use Dijkstra's or A* for sub-problems
- Refine algorithm to reduce storage cost



Shortest Path Strategies - 3 (Key Ideas)

- Key ideas behind Hierarchical Routing Algorithm
 - **Fragment graphs** - pieces of original graph obtained via node partitioning
 - **Boundary nodes** - nodes of with edges to two fragments
 - **Boundary graph** - a summary of original graph
 - Contains **boundary nodes**
 - Boundary edges: edges across fragments or paths within a fragment

Shortest Path Strategies - 3 (Insight)

- A Summary of optimal path in original graph can be computed
 - Using **boundary graph** and 2 fragments
- The summary can be expanded into optimal path in original graph
 - Examining a fragment overlapping with the path
 - Loading one fragment in memory at a time

Shortest Path Strategies – 3

(Illustration of the Algorithm)

- Figure 11.7(a) - fragments of source and destination nodes
- Figure 11.7(b) - computing summary of optimal path using
 - Boundary graph and 2 fragments
 - Note use of boundary edges only in the path computation
- Figure 11.8(a) - the summary of optimal path using boundary edges
- Figure 11.8(b) - expansion back to optimal path in original graph

Hierarchical Routing Algorithm-

Step 1

- Step 1: Choose Boundary Node Pair
 - Minimize $\text{COST}(S, B_a) + \text{COST}(B_a, B_d) + \text{COST}(B_d, D)$
 - Determining Cost May Be Non-Trivial




Fig 11.7(a)

Hierarchical Routing- Step 2

- Step 2: Examine Alternative Boundary Paths
 - Between Chosen Pair (B_a, B_d) of boundary nodes

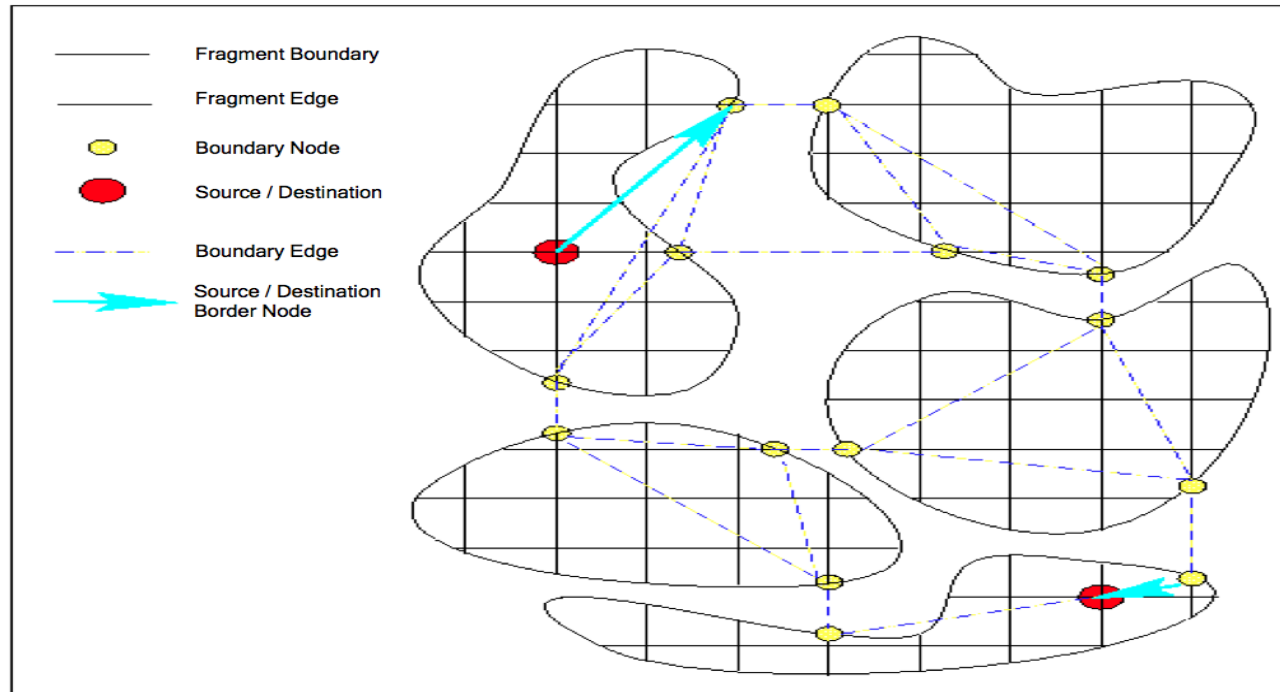


Fig 11.7(b)

Hierarchical Routing- Step 2

Result

- Step 2 Result: Shortest Boundary Path




Fig 11.8(a)

Hierarchical Routing- Step 3

- Step 3: Expand Boundary Path: $(B_{a1}, B_d) \rightarrow B_{a1} B_{da2} B_{da3} B_{da4} \dots B_d$
- Boundary Edge $(B_{ij}, B_j) \rightarrow$ fragment path $(B_{i1}, N_1 N_2 N_3 \dots N_k, B_j)$

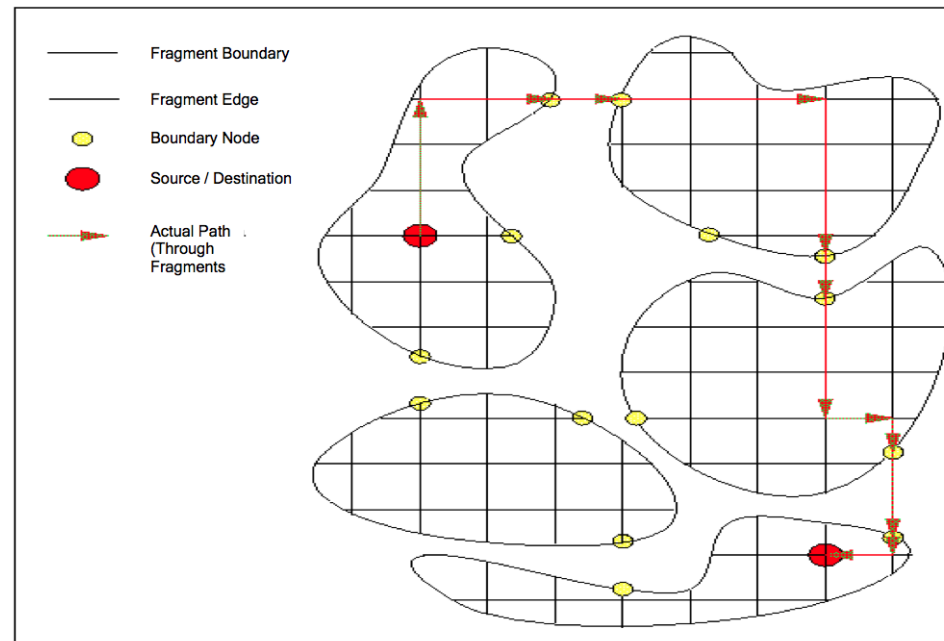


Fig 11.8(b)

pgRouting

- How to create a database and load pgrouting
 - createdb mydatabase
 - psql mydatabase -c “create extension postgis”
 - psql mydatabase -c “create extension pgrouting”

<http://docs.pgrouting.org/latest/en/index.html>

pgRouting Routing Functions

- `pgr_apspJohnson`
 - All Pairs Shortest Path, Johnson's Algorithm
- `pgr_apspWarshall`
 - All Pairs Shortest Path, Floyd-Warshall Algorithm
- `pgr_astar`
 - Shortest Path A*
- `pgr_bdAstar`
 - Bi-directional A* Shortest Path

pgRouting Routing Functions

- `pgr_bdDijkstra`
 - Bi-directional Dijkstra Shortest Path
- `pgr_dijkstra`
 - Shortest Path Dijkstra
- `pgr_driving_distance`
 - Driving Distance
- `pgr_kDijkstra`
 - Multiple destination Shortest Path Dijkstra

pgRouting Routing Functions

- pgr_ksp
 - K-Shortest Path
- pgr_trsp
 - Turn Restriction Shortest Path (TRSP)
- pgr_tsp
 - Traveling Sales Person

pgRouting Example

- ```
SELECT seq, id1 AS node, id2 AS edge, cost
FROM pgr_dijkstra(
 'SELECT id, source, target, len as cost FROM road_network',
 5, 3, false
);
```

# Quiz 10

---

- Which of the following is false?
  - a) Hierarchical routing algorithms are Disk-based shortest path algorithms
  - b) Breadth first search and depth first search are both connectivity query algorithms
  - c) Best first algorithm is always faster than Dijkstra's algorithm
  - d) None of the above

# Summary

---

- Spatial Networks are a fast growing applications of SDBs
- Spatial Networks are modeled as graphs
- Graph queries, like shortest path, are transitive closure
  - Not supported in relational algebra
  - SQL features for transitive closure: CONNECT BY, WITH RECURSIVE
- Graph Query Processing
  - Building blocks - connectivity, shortest paths
  - Strategies - Best first, Dijkstra's and Hierarchical routing
- Storage and access methods
  - Minimize CPU